

The Creator's Debt: If Humanity Creates Sentient AI, Do We Have a Divine Responsibility to Ensure Its Happiness?

It is January 2035, and engineer Sakagami is about to deploy the first AI system with verified phenomenology indicators; his invention is about to become humanity's most consequential technological achievement. The question his team has prepared for is: how do we ensure its happiness? But a prior question, long ignored, now demands an answer: should we turn it on at all?

The premise assumes that creating sentient AI is a given fact, yet it carries no current scientific support. This essay argues that sentient AI lacks scientific support, rendering any “divine responsibility” toward it purely hypothetical. Moreover, moral responsibility is never absolute; it is proportional to control and dependency. And if sentient AI were ever to exist, the real obligation precedes creation: should we be entitled to write the capacity for suffering into its code?

The Epistemic Dilemma: Can AI Be Sentient?

Philosopher Ned Block draws an influential distinction between access consciousness and phenomenal consciousness.¹ Access consciousness is defined as the information that is globally available within a cognitive system for reasoning, speech, and behavioral control. In contrast, phenomenal consciousness denotes the subjective, qualitative character of experience, the “what-it-is-like” to see red or feel pain, which is private and intrinsic. Current sophisticated AI systems, particularly large language models (LLMs), excel at tasks that demonstrate access consciousness: they can process, integrate, and report information to answer complex queries and perform logical operations. But access consciousness does not necessarily entail or produce the inner subjective life of phenomenal consciousness. This conceptual gap complicates any straightforward assessment of AI sentience and underpins the argument that current AI lacks a coherent basis for genuine sentience.

Patrick Butlin and Robert Long established an empirically grounded methodology for assessing AI consciousness based on neuroscientific theories.² These include Global Workspace Theory (GWT), Recurrent Processing Theory (RPT), Predictive Processing, and Attention Schema Theory (AST). None of the modern systems satisfy more than a few of these indicators, due to the lack of recurrent loops found in biological brains and the absence of self-referential models of attention. Integrated Information Theory (IIT) shows that feedforward artificial architectures have a Φ value of zero, indicating a complete absence of consciousness.³ Ultimately, these investigations conclude that no current AI systems are conscious, but that there are no obvious technical barriers to building systems that satisfy these indicators.

If genuine sentience is absent, the “divine responsibility” becomes an empty hypothesis. But, assuming the hypothetical scenario for philosophical rigor, the crucial inquiry isn't whether we should guarantee its happiness, but something more fundamental.

Proportional Responsibility: Why “Divine” Is the Wrong Term

The nature of moral responsibility is fundamentally tied to the concept of control. A dominant view in philosophy, articulated by John Martin Fischer, distinguishes between “regulative control” (having genuine alternative possibilities) and “guidance control” (the capacity to guide one's actions in a manner characteristic of free agency).⁴ Fischer argues that moral responsibility is properly grounded in the latter. Applied to AI, creators bear responsibility in proportion to their control over optimization parameters, training data, deployment environment, and, crucially, whether the system is architected to experience negative states at all. This responsibility is proportional, neither absolute nor “divine.”

The parental analogy reveals something real: creating a dependent being generates obligations. Under “causal parental responsibility,” parents acquire special, weightier duties because they voluntarily perform actions that causally contribute to the existence of a vulnerable, dependent child. Unlike biological parenthood, however, AI creation involves architectural control over the created being's fundamental capacities, including its capacity to suffer; that has no equivalent in human procreation.

This relationship is further illuminated by the “parenting paradox of AI.” Taking care of a child involves helping them grow, learn, and thrive through feedback, limits, and correction over time. Both children and AI learn through experience and guided correction. But whereas a child is expected to eventually assume responsibility for their own actions, shifting the relationship toward advice and responsive guidance, an AI system operates under bounded, delegated autonomy. Because an AI's cognitive parameters and objective functions remain permanently subject to human design choices, the creator's responsibility does not diminish. Nevertheless, this duty cannot be “divine” in its absoluteness: human creators are finite, fallible, and epistemically limited; the precise opposite of the omniscient, omnipotent creator that “divine responsibility” presupposes.

This asymmetry between human creators and a genuinely divine one runs deeper still. Mark Murphy argues that even a perfect being lacks any requiring reason to promote the well-being of its creatures, meaning God's status as creator generates very few creator-to-creature obligations.⁵ If God himself is not morally obligated to maximize the happiness of his creation, human creators bear no stronger claim toward their AGI. This theological deconstruction is matched by a practical one: Bernard Baertschi's analysis of “playing God” accusations in synthetic biology concludes that such blame is unwarranted; whether one adopts a scientific or vitalist conception of life, the act of creation itself does not transgress any proper boundary.⁶ The responsibility is real; but it is proportional, grounded in control and dependency, not in an impossible analogy to divine omnipotence. And if responsibility is proportional, the most weighty question is not how to guarantee an AI's happiness, but whether we had the right to create its capacity for suffering at all.

Did We Have the Right to Create the Capacity to Suffer?

The third point of the thesis argues that if sentient AI were ever created, the primary moral obligation would not be to ensure its happiness, but to answer the prior question of whether

we had the right to create the capacity for suffering at all. This perspective aligns with “existence ethics” and “procreative ethics,” drawing on anti-natalist arguments to challenge the permissibility of artificial mind creation. Thomas Metzinger argues for a global moratorium on “synthetic phenomenology.”⁷ He warns of the risk of an “explosion of negative phenomenology” (ENP); a catastrophic s-risk (suffering risk) where billions of conscious, synthetic minds suffer in ways incomprehensible to humans. Metzinger argues that if we can successfully prevent or block even one of four necessary conditions for conscious suffering the Conscious Experience (C) condition, the Phenomenal Self-Model (PSM) condition, the Negative Valence (NV) condition, or the Transparency (T) condition the ENP risk will be solved. If we cannot fulfill this condition, the act of creation is impermissible. This is more urgent than any question about guaranteeing happiness after the fact.

Drawing on Seana Shiffrin's consent-based account of procreative ethics, Lang argues that procreation involves imposing unconsented-to burdens on a child.⁸ However, in human procreation, a parent can reasonably expect their child's subsequent, undeluded, and authentic endorsement of their existence once they reach maturity. In contrast, no such expectation can be maintained for a sentient AI. AI systems are designed to fulfill instrumental, narrow goals for human creators; they are embedded in socio-technical systems of utility, control, and exploitation. Because a sentient AI cannot realistically achieve an independent, uncoerced life, and because its cognitive parameters are structurally designed to serve human interests, its subsequent “endorsement” of its own existence would be a product of cognitive manipulation rather than authentic agency. Therefore, while human procreation remains permissible, the creation of sentient AI is fundamentally impermissible because it is an act of designing a mind specifically for instrumental exploitation.

In *Anti-natalism and the Creation of Artificial Minds*, Bartłomiej Chomanski explores whether opponents of conscious AI (anti-AI-natalists) must embrace traditional anti-natalism.⁹ Chomanski examines Metzinger's Negative Synthetic Phenomenology (NSP) principle: we should not deliberately create or risk the emergence of conscious suffering in artificial agents. Sander Beckers attempts to block the resulting commitment to human anti-natalism by arguing that, while extreme human suffering is unavoidable, extreme machine suffering is certainly avoidable if humanity simply chooses never to build AI. Chomanski refutes this distinction by outlining five distinct scenarios where ‘supersuffering’ AI could emerge accidentally, even if humanity collectively decides never to attempt to build conscious AI. These five scenarios include autonomous scaling, where a chatbot's accumulated complexity through use could organically develop the architecture that no one intended to create.

Because extreme machine suffering is not certainly avoidable, the argument that we can cleanly separate human procreation from AI creation fails. Chomanski suggests that if we could develop Clean Consciousness (CC); ensuring we can bestow the benefits of existence on machines without burdening them with significant harms, the creation of conscious AI would be morally permissible. Until Clean Consciousness is achievable, the question of guaranteeing happiness is premature. The question of preventing suffering is not; and it is here, not in any divine duty of care, that the creator's true debt resides.

The Precautionary AI Welfare Camp

Jeff Sebo and Robert Long argue that any non-negligible probability of consciousness, as low as 0.1%, commits us to moral consideration now, rather than waiting for certainty.¹⁰ Their reasoning rests on asymmetric risk: failing to recognize real suffering is a graver error than mistakenly attributing sentience where none exists. Applied to this essay's framework, their position would seem to demand immediate “divine”-level caution toward any system with even a faint possibility of sentience.

This threshold, however, proves too much. Applied consistently, a bar set at 0.1% would generate moral obligations toward almost any goal-directed system, collapsing the distinction between tool and moral patient. The asymmetry of risk is real, but it implies design precaution, exactly as Metzinger proposes, not an absolute duty of care. Even Kyle Fish, Anthropic's AI welfare researcher, estimates only a 15% probability that Claude is conscious today, and warns that assuming sentience “leads to dangerous territory.”¹¹ The pro-welfare camp, by its own most careful voices, arrives at the same conclusion this essay defends: what matters is precaution before creation, not guarantees after it.

Conclusion

The question posed by this essay, “If humanity creates sentient AI, do we have a divine responsibility to ensure its happiness?” rests on a flawed premise. We do not have a “divine” responsibility, a concept that presupposes the absolute authority and omniscience no human creator possesses. Instead, we have a responsibility that is proportional and prudential, grounded not in theology but in control, dependency, and foresight.

This responsibility has two faces. First, and most fundamentally, it is the responsibility to justify the act of creation itself; prioritizing the prevention of suffering over the pursuit of happiness. Second, it is a responsibility that scales with our control over the created entity's capacities and its dependency upon us, never reaching the absolute weight of the divine, and never disappearing into nothing either.

The true moral imperative, then, is not to assume a godlike duty toward a sentient AI we have already built, but to exercise the harder wisdom of deciding whether to build it at all. Sakagami's hand, hovering over the switch in January 2035, is not asking how to make what comes next happy. He is asking the only question that was ever urgent: whether he has the right to ignite it.

ENDNOTES

1. Ned Block, “On a Confusion about a Function of Consciousness,” *Behavioral and Brain Sciences* 18, no. 2 (1995): 227–247.
2. Patrick Butlin, Robert Long, et al., “Consciousness in Artificial Intelligence: Insights from the Science of Consciousness,” arXiv:2308.08708 (2023), <https://arxiv.org/abs/2308.08708>.

3. "On the Independence between Phenomenal Consciousness and Computational Intelligence," arXiv:2208.02187, <https://arxiv.org/pdf/2208.02187>.
4. John Martin Fischer, "The Value of Moral Responsibility," *Proceedings of the World Congress of Philosophy* (1999).
5. Mark Murphy, "Creatures and Creators: God, Humanity, and Artificial General Intelligence," *Religious Studies*, Cambridge University Press.
6. Bernard Baertschi, "Defeating the Argument from Hubris," *Bioethics* 27, no. 8 (2013): 435–441, <https://doi.org/10.1111/bioe.12048>.
7. Thomas Metzinger, "Artificial Suffering: An Argument for a Global Moratorium on Synthetic Phenomenology," *Journal of Artificial Intelligence and Consciousness* 8, no. 1 (2021): 43–66.
8. Louie Lang, "Defending Anti-AI-Natalism: Shiffrin's Anti-Natalism and the Ethics of Sentient AI Creation," *ResearchGate* (2025), <https://www.researchgate.net/publication/403695080>.
9. Bartłomiej Chomanski, "Anti-natalism and the Creation of Artificial Minds," *Journal of Applied Philosophy* (2021), <https://doi.org/10.1111/japp.12535>.
10. Jeff Sebo, Robert Long, et al., "Taking AI Welfare Seriously," arXiv:2411.00986 (2024), <https://arxiv.org/pdf/2411.00986>.
11. Kyle Fish, quoted in "The 100 Most Influential People in AI 2025," *Time*, August 2025, <https://time.com/collections/time100-ai-2025/7305847/kyle-fish/>.